

Data Science Nigeria: Introductory Machine Learning Training



WEEK 4: INTRODUCTION TO CLASSIFICATION

Binary/ Multi-class classification

Classification Algorithms

Evaluation of Classification Models Performance

Error/ Cost function

Confusion matrix

Precision

AUC

In Classification, we predict the category a data belongs to ie. Classification algorithms are used to predict labels

- Spam Detection
- Churn Prediction
- Sentiment Analysis
- Dog Breed Detection

TYPES OF CLASSIFICATION TASK

- Binary classification eg. e-mail spam detection (1 ->spam; or 0→not spam), biometric identification, whether a customer will default or Not
- Multi-class classification eg. digit recognition (where classes go from 0 to 9), predicting a party that wins the election,

Classification Algorithms

1. Logistic Regression
2. Naive Bayes Classifier
3. Nearest Neighbor
4. Support Vector Machines
5. Decision Trees
6. Boosted Trees
7. Random Forest

Import Modules

In [148]:

```
#import modules

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

%matplotlib inline
```

In [149]:

```
data= pd.read_csv("Social_Network_Ads.csv", delimiter= ",")
data.head()
```

Out[149]:

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

In [150]:

```
data['Gender'].value_counts()
```

Out[150]:

```
Female    204
Male      196
Name: Gender, dtype: int64
```

In []:

Pre-processing Data

In [151]:

```
#creating a LabelEncoder object
from sklearn.preprocessing import LabelEncoder

le= LabelEncoder()
#invoking fit_transform method on object
data['Gender']=le.fit_transform(data['Gender'])
```

In [152]:

```
data['Gender'].value_counts()
```

Out[152]:

```
0    204
1    196
Name: Gender, dtype: int64
```

Exploratory analysis

In [153]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
User ID      400 non-null int64
Gender       400 non-null int32
Age          400 non-null int64
EstimatedSalary  400 non-null int64
Purchased    400 non-null int64
dtypes: int32(1), int64(4)
memory usage: 14.1 KB
```

In [154]:

```
# employees that did not buy and those that bought

left= data.groupby('Purchased')
left.mean()
```

Out[154]:

	User ID	Gender	Age	EstimatedSalary
Purchased				
0	1.569116e+07	0.505837	32.793774	60544.747082
1	1.569222e+07	0.461538	46.391608	86272.727273

In [155]:

```
data.describe()
```

Out[155]:

	User ID	Gender	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000	400.000000
mean	1.569154e+07	0.490000	37.655000	69742.500000	0.357500
std	7.165832e+04	0.500526	10.482877	34096.960282	0.479864
min	1.556669e+07	0.000000	18.000000	15000.000000	0.000000
25%	1.562676e+07	0.000000	29.750000	43000.000000	0.000000
50%	1.569434e+07	0.000000	37.000000	70000.000000	0.000000
75%	1.575036e+07	1.000000	46.000000	88000.000000	1.000000
max	1.581524e+07	1.000000	60.000000	150000.000000	1.000000

Data Visualization

Users that purchased these Ads

In [188]:

```
left_count = left.count()
```

In [189]:

```
left_count
```

Out[189]:


	User ID	Gender	Age	EstimatedSalary
Purchased				
0	257	257	257	257
1	143	143	143	143

In [156]:



```
#how many users were in each category ?
```

```
plt.bar(left_count.index.values, left_count['User ID'])  
plt.ylabel("Numbers of Users")  
plt.show()
```



In [161]:



```
# ratio of users that did bought the Ads were only 36%
```

```
data.Purchased.value_counts()  
float(data.Purchased.value_counts()[1])/len(data) * 100
```

Out[161]:

35.75

A little bit of feature engineering !

1) Using an Age category

IGen[1-24], Millenials[24-39], GenX[40-54], BabyBoomers[55-73]

In [162]:

```
data['Age Category'] = ['iGen' if 0<age<25 else 'Millenials' if 24<age<40 else 'GenX' if 40<=age<65 else 'Boomers' if 65<=age]
for age in list(data['Age'].values)
```

In [163]:

```
data.head()
```

Out[163]:

	User ID	Gender	Age	EstimatedSalary	Purchased	Age Category
0	15624510	1	19	19000	0	iGen
1	15810944	1	35	20000	0	Millenials
2	15668575	0	26	43000	0	Millenials
3	15603246	0	27	57000	0	Millenials
4	15804002	1	19	76000	0	iGen

2) Using an Income category [Inter-quartile Ranges]

In [164]:

```
data['Income Category'] = pd.qcut(data['EstimatedSalary'],3,labels=['Low','Medium','High'])
```

In [165]:

```
data.head()
```

Out[165]:

	User ID	Gender	Age	EstimatedSalary	Purchased	Age Category	Income Category
0	15624510	1	19	19000	0	iGen	Low
1	15810944	1	35	20000	0	Millenials	Low
2	15668575	0	26	43000	0	Millenials	Low
3	15603246	0	27	57000	0	Millenials	Medium
4	15804002	1	19	76000	0	iGen	Medium

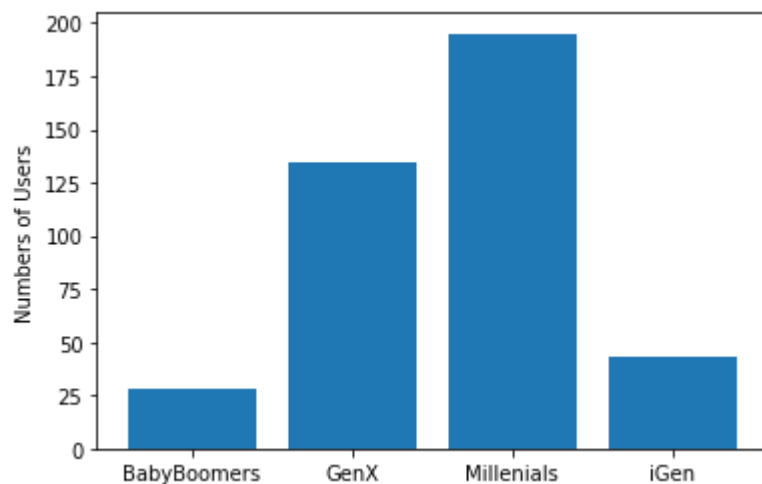
Lets see the Visuals

In [166]:

```
age_cat = data.groupby('Age Category').count()
inc_cat = data.groupby('Income Category').count()
```


In [167]:

```
plt.bar(age_cat.index.values, age_cat['Age'])  
plt.ylabel("Numbers of Users")  
plt.show()
```



In [168]:

```
plt.bar(inc_cat.index.values, inc_cat['Age'])  
plt.ylabel("Numbers of Users")  
plt.show()
```



In []:

```
# data[data['Income Category'] == 'High'].max()
```

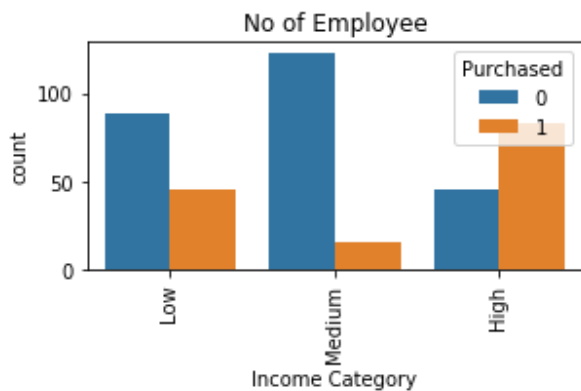
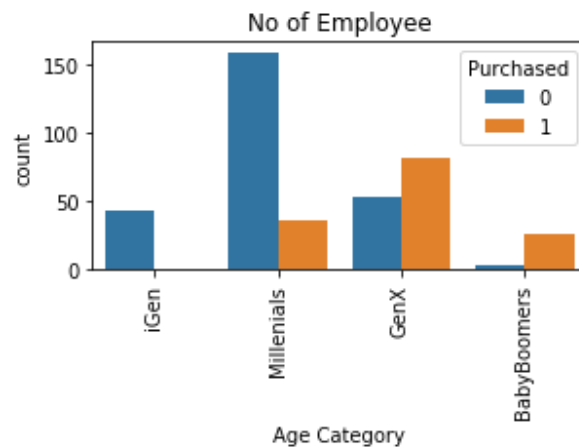

In [169]:



```
#whats the statistics based on those who did not make purchase
```

```
features= ['Gender', 'Age Category', 'Income Category']  
fig= plt.subplots(figsize= (10,15))
```

```
for i,j in enumerate(features):  
    plt.subplot(4,2, i+1)  
    plt.subplots_adjust(hspace=1.0)  
    sns.countplot(x=j, data=data , hue= "Purchased")  
    plt.xticks(rotation= 90)  
    plt.title("No of Employee")
```



Model building

In [171]:

```
data.columns
```

Out[171]:

```
Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased',  
      'Age Category', 'Income Category'],  
      dtype='object')
```

In [172]:

```
# split data into features and target  
  
x = data[['Gender', 'Age', 'EstimatedSalary']]  
y=data['Purchased']  
  
#train_test_split  
from sklearn.model_selection import train_test_split  
  
x_train, x_test,y_train, y_test= train_test_split (x,y,test_size=0.3,random_state=42)
```

Assignment: Use engineered features i.e new columns in training your model. Confirm if it improves our models significantly or not ?

Algorithm/ Model 1 : Naive Bayes

In [173]:

```
# Fitting Naive Bayes to the Training set  
from sklearn.naive_bayes import GaussianNB  
classifier = GaussianNB()  
classifier.fit(x_train, y_train)
```

Out[173]:

```
GaussianNB(priors=None, var_smoothing=1e-09)
```

In [174]:

```
# Predicting the Test set results  
y_pred = classifier.predict(x_test)
```

In [175]:

```
# evaluating performance : Accuracy, Precision, Recall  
from sklearn import metrics  
  
print("Accuracy: ", metrics.accuracy_score(y_test, y_pred))  
print("Precision: ", metrics.precision_score(y_test, y_pred))  
print("Recall: ", metrics.recall_score(y_test, y_pred))
```

```
Accuracy:  0.925  
Precision: 0.975  
Recall:    0.8297872340425532
```

In [176]:

```
# calculate ROC Curve
from sklearn.metrics import roc_curve, auc
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)
print ("ROC AUC", roc_auc)
```

ROC AUC 0.9080443019527833


In [177]:

```
# Plot confusion Matrix
def conftable(test, pred, imagename):
    confmatrix = metrics.confusion_matrix(y_test, y_pred)
    plt.matshow(confmatrix)
    plt.title('Confusion matrix')
    plt.colorbar()
    plt.ylabel('GroundTruth')
    plt.xlabel('Predicted')
    plt.savefig(imagename)

    plt.show()
    print(confmatrix)
```

In [178]:

```
conftable(y_test, y_pred, "conf")
```



```
[[72  1]
 [ 8 39]]
```

In [179]:

```
# Ground Truth
pd.Series(y_test).value_counts()
```

Out[179]:

```
0    73
1    47
Name: Purchased, dtype: int64
```

Algorithm/ Model 2 : Logistic Regression

In [180]:



```
# Fitting Logistic Regression to the Training set
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state = 0)
classifier.fit(x_train, y_train)
```

```
C:\Users\training.NG-1NW8PX1\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.
  FutureWarning)
```

Out[180]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                  intercept_scaling=1, max_iter=100, multi_class='warn',
                  n_jobs=None, penalty='l2', random_state=0, solver='warn',
                  tol=0.0001, verbose=0, warm_start=False)
```

In [181]:




```
# Predicting the Test set results
y_pred = classifier.predict(x_test)
```

In [182]:

```
print("Accuracy: ", metrics.accuracy_score(y_test, y_pred))
print("Precision: ", metrics.precision_score(y_test, y_pred))
print("Recall: ", metrics.recall_score(y_test, y_pred))
fpr, tpr, thresholds = roc_curve (y_test, y_pred)
roc_auc= auc (fpr, tpr)
print ("ROC AUC", roc_auc)
confable(y_test,y_pred,"conf")
```

Accuracy: 0.6083333333333333
Precision: 0.0
Recall: 0.0
ROC AUC 0.5

C:\Users\training.NG-1NW8PX1\Anaconda3\lib\site-packages\sklearn\metrics\classification.py:1143: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples.
'precision', 'predicted', average, warn_for)



```
[[73  0]
 [47  0]]
```

Algorithm/ Model 3 : Random Forest

In [183]:

```
# Fitting Random Forest Classification to the Training set
from sklearn.ensemble import RandomForestClassifier
classifier = RandomForestClassifier(n_estimators = 10, criterion = 'entropy', random_state=0)
classifier.fit(x_train, y_train)
```

Out[183]:

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='entropy',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                        oob_score=False, random_state=0, verbose=0, warm_start=False)
```


In [184]:

```
# Predicting the Test set results
y_pred = classifier.predict(x_test)
```

In [185]:

```
print("Accuracy: ", metrics.accuracy_score(y_test, y_pred))
print("Precision: ", metrics.precision_score(y_test, y_pred))
print("Recall: ", metrics.recall_score(y_test, y_pred))
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)
print("ROC AUC", roc_auc)
confable(y_test, y_pred, "conf")
```

Accuracy: 0.9
Precision: 0.8888888888888888
Recall: 0.851063829787234
ROC AUC 0.8912853395511511



```
[[68  5]
 [ 7 40]]
```

In []:

Extra Algorithm/ Model 4 : XGBoost

In [186]:

```
import xgboost as xgb
from xgboost import XGBClassifier


xgboost = XGBClassifier()
xgb = xgboost.fit(x_train, y_train)
y_pred = xgb.predict(x_test)
```

In [187]:



```
print("Accuracy: ", metrics.accuracy_score(y_test, y_pred))
print("Precision: ", metrics.precision_score(y_test, y_pred))
print("Recall: ", metrics.recall_score(y_test, y_pred))
fpr, tpr, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(fpr, tpr)
print("ROC AUC", roc_auc)
confable(y_test, y_pred, "conf")
```

Accuracy: 0.9333333333333333
Precision: 0.9148936170212766
Recall: 0.9148936170212766
ROC AUC 0.9300495482366656



```
[[69  4]
 [ 4 43]]
```

Assignment

Extract other features and re-train the classification models to note their accuracy metrics